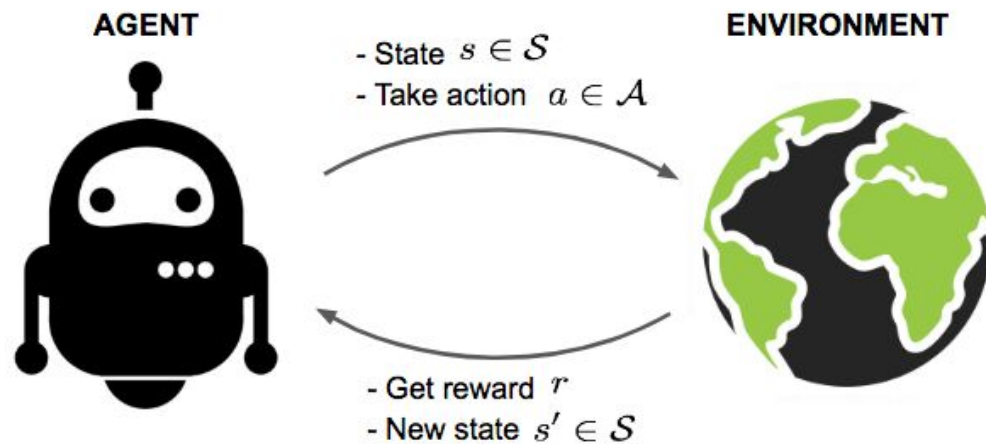


Ensemble Learning

Teaching Agents to Walk

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

Problem Statement



Problem Statement – Contd.

- Discounted returns represent the sum of all rewards ' r_t ' ever obtained by the agent discounted by $\gamma^t \in (0,1)$

$$R_t = \sum_{t=0}^{\infty} \gamma^t r_t$$

- Goal of RL is to learn the optimal policy ' π_{θ} ' which maximizes the expected return:

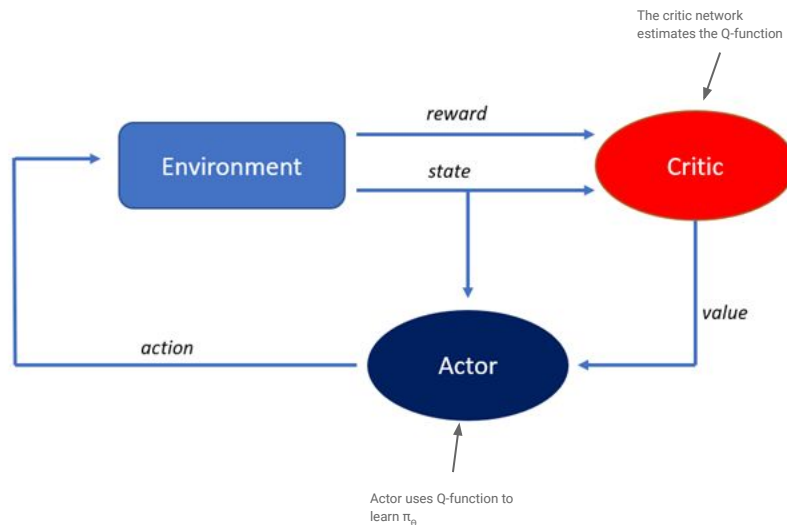
$$\pi_{\theta}^* = \max_{\pi_{\theta}} \mathbb{E}_{\pi_{\theta}} [R_t]$$

Background

- Our proposed algorithm builds upon TD3
- We will go over the necessary information:
 - ◆ Deep Deterministic Policy Gradients (DDPG)
 - ◆ Twin Delayed DDPG (TD3)

Deep Deterministic Policy Gradients (DDPG)

- Based on the actor-critic framework:
 - ◆ **Actor Network:** Learns the deterministic policy (π_{θ})
 - ◆ **Critic Network:** Approximates the Q-function
- DDPG is a model-free, off-policy algorithm.
- Makes use of two target networks:
 - ◆ Target Actor Network
 - ◆ Target Critic Network



DDPG: Q-Function Estimation

→ The target is approximated using Temporal Difference in conjunction with the secondary target networks:

The diagram shows the equation $y = r(s, a) + \gamma Q_{\phi'}(s', a')$ with four annotations and arrows pointing to specific parts of the equation:

- An arrow from the label "Target" points to the variable y .
- An arrow from the label "Present rewards" points to the term $r(s, a)$.
- An arrow from the label "Target Network" points to the function $Q_{\phi'}$.
- An arrow from the label " a' sampled from target actor" points to the variable a' .

→ Using the target, we minimize the following error:

$$L(\phi') = \mathbb{E}_{s,a,s',a'} [(Q_{\phi}(s, a) - y)^2]$$

DDPG: Policy Eval & Replay Buffer

- Policy Evaluation:

- Since the Q-function is differentiable, we can just perform gradient ascent to find the optimal policy:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{s \sim p_{\pi}} [\nabla_a Q_{\pi}(s, a)|_{a=\pi_{\theta}(s)} \nabla_{\theta} \pi_{\theta}(s)]$$

- Experience Replay Buffer:


- DDPG uses an experience replay buffer to store previous experiences (states, action, reward).
- Prevents bias since batches of samples are drawn from it at random.

Twin Delayed DDPG (TD3)

- Successor to DDPG
- Addresses the overestimation of Q-values faced by DDPG
- Mitigates the problem by three tasks:
 - ◆ **Target Policy Smoothing:** Adds noise to the target actions, \mathbf{a}'
 - ◆ **Clipped Q-Learning:** Learns two Q-function
 - ◆ **Delayed Policy Updates:** Policy and target policy are updated after every other step

Multi-TD3


- Based on the TD3 algorithm
- Works on the principle of “Ensemble Learning”
- Performs the following tasks:
 - ◆ Two separate instances of TD3 networks are trained
 - ◆ Both networks push and sample from the same replay buffer
 - ◆ Model with highest rewards serves as the model to be trained in a new episode by the two networks
- Advantage:
 - ◆ Decreases the likelihood of selecting a relatively poor model
 - ◆ More experiences are sampled per episode




Multi-TD3: Ensemble Learning

Results


All trainings were conducted using OpenAI's 3 gym environments



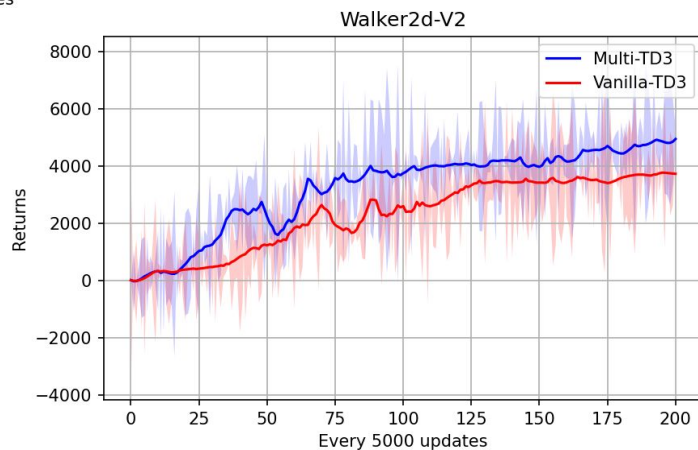
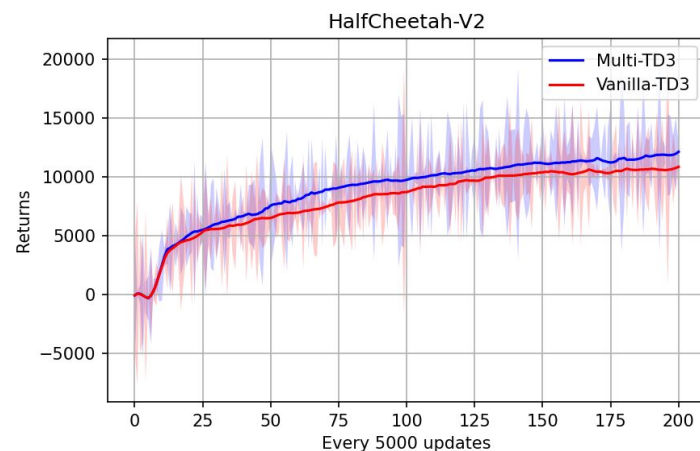
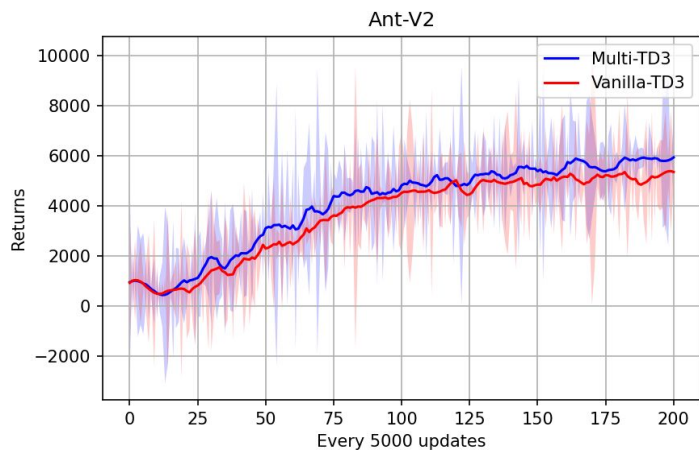
Walker2d-v2





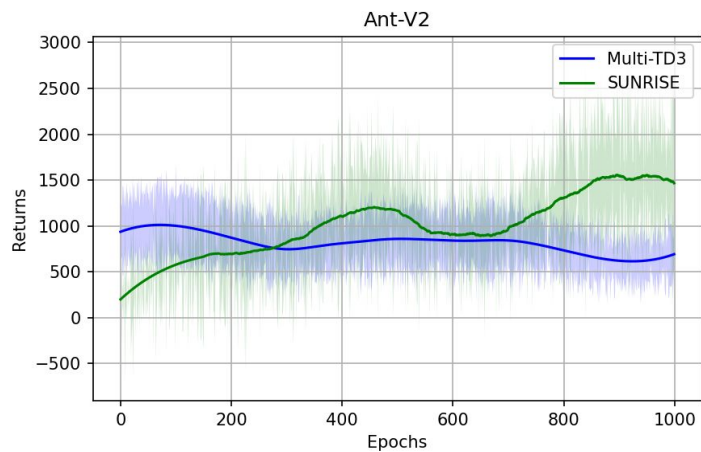
HalfCheetah-v2



Ant-v2



Results: Multi-TD3 vs TD3



Results: Multi-TD3 vs SUNRISE

Discussion

Limitations

- Not enough time to compare TD3, Multi-TD3, and SUNRISE on the same timescale
- Because of long training times, no hyperparameter tuning could be done
 - ◆ the size of the replay buffer
 - ◆ the number of past trajectories sampled

Applications

- TD3 has found numerous applications in robotics, for example teaching a robotic arm to reach a target; Multi-TD3 should improve it.
- We already saw the improvement in the learned policy for Walker2d-V2

Thank you!

The End